

Kanishk Jain

📱 mobile | @ email | in linkedin | github | 🌐 webpage

EDUCATION

Université de Montréal | Mila

Montréal, Canada

PhD in Computer Science and Engineering; GPA: 4.3

Sep 2023 – Present

- Research Focus: Reinforcement Learning, Multi-Agent Frameworks, Foundational Multimodal Models, Lifelong Learning, Multimodal Representation Learning
- Advisors: Prof. Sarath Chandar, Prof. Ross Goroshin

International Institute of Information Technology

Hyderabad, India

MS by Research in Computer Science and Engineering; GPA: 9.33/10

Aug 2021 – Dec 2022

- Research Focus: Multimodal Learning, Visual Grounding, Language-Guided Autonomous Navigation, Semi-Supervised Learning
- Advisor: Dr. Vineet Gandhi

RESEARCH PUBLICATIONS

Discovering Failure Modes in Vision-Language Models using RL

Paper

Kanishk Jain, Qian Yang, Shravan Nayak, Parisa Kordjamshidi, Nishanth Anand, Aishwarya Agrawal *arXiv 2026*

- An RL framework that trains a questioner agent to generate increasingly complex queries, automatically exposing weaknesses in vision-language models without human intervention.

Prioritized Concept Learning via Relative Error-driven Sample Selection

Paper

Shivam Chandhok, Qian Yang, Oscar Manas, **Kanishk Jain**, Leonid Sigal, Aishwarya Agrawal

CVPR 2026

- A data-efficient training framework for VLMs that dynamically selects samples based on model's learning propensity

Benchmarking Vision Language Models for Cultural Understanding

Paper

Shravan Nayak, **Kanishk Jain**, Rabiul Awal, Siva Reddy, et al.

EMNLP 2024 (Oral)

- Introduced CulturalVQA, a benchmark for evaluating Vision Language Model's understanding of geo-diverse cultural elements, revealing performance disparities across regions and cultural facets.

Test-Time Amendment with a Coarse Classifier for Fine-Grained Classification

Paper

Kanishk Jain, Shyamgopal Karthik, Vineet Gandhi

NeurIPS 2023

- Proposed Hierarchical Ensembles (HiE), a post-hoc correction strategy that utilizes coarse-grained predictions at test-time to significantly reduce mistake severity and enhance fine-grained classification accuracy

Instance-Level Semantic Maps for Vision Language Navigation

Paper

Laksh Nanwani, Anmol Agarwal, **Kanishk Jain**, et al.

ROMAN 2023

- Developed Instance-Level Semantic Maps to resolve object ambiguity, enabling navigation agents to precisely ground complex natural language instructions.

Ground then Navigate: Language-guided Navigation in Dynamic Scenes

Paper

Kanishk Jain*, Varun Chhangani*, Amogh Tiwari, K. Madhava Krishna, Vineet Gandhi

ICRA 2023

- Developed a visual-grounding-based planning framework that maps language instructions to navigable road regions, enabling navigation in dynamic outdoor environments.

Bringing Generalization to Deep Multi-View Pedestrian Detection

Paper

Jeet Vora, Swetanjal Dutta, **Kanishk Jain**, Shyamgopal Karthik, Vineet Gandhi

WACV-W 2023

- Developed an evaluation framework and synthetic dataset to benchmark generalization in multi-view pedestrian detection to unseen configurations and realistic scenarios.

Comprehensive Multi-Modal Interactions for Referring Image Segmentation

Paper

Kanishk Jain, Vineet Gandhi

ACL Findings 2022

- Developed a Transformer-based architecture for Referring Image Segmentation, leveraging synchronous multi-modal attention and hierarchical aggregation to enhance segmentation performance.

Grounding Linguistic Commands to Navigable Regions

Paper

Kanishk Jain*, Nivedita Rufus*, Unni Krishnan*, Vineet Gandhi, K. Madhava Krishna

IROS 2021

- Introduced Referring Navigable Regions (RNR), a novel task for grounding linguistic commands to navigable road regions, enabling autonomous agents to follow human instructions for navigational maneuvers.

WORK EXPERIENCE

Noah's Ark Lab | Huawei

Intern Researcher

Sep 2025 – Jan 2026

- Developing modality-invariant foundation models to mitigate modality competition by leveraging contrastive learning for robust multimodal representations.

Université de Montréal | IFT 6765

Teaching Assistant

Jan 2025 – Apr 2025

- Held weekly office hours, designed and graded assignments, and mentored teams on project ideas.

Université de Montréal | Mila

PhD Student

Sep 2023 – Present

- Developing a Reinforcement Learning framework that automatically discovers failure modes in Vision-Language Models, training an adversary to generate inputs that expose systematic blind spots and produce targeted data for robustness in safety-critical applications.
- Exploring early-exit methods for streaming video understanding to enable real-time decision making with reliability guarantees.

CVIT, IIIT Hyderabad

Research Engineer

Mar 2023 – Aug 2023

- Developed Hierarchical Ensembles to reduce mistake severity in fine-grained classification, improving performance in supervised and semi-supervised settings.
- Designed and deployed a real-time player tracking solution in Bird's Eye View, currently utilized in live cricket broadcasts.

SELECTED PROJECTS

T2I Generation using SSMs: Analyzing the effectiveness of Selective State Space Models (SSM) in reducing computational complexity for text-to-image generation.

Top-View Player Tracking: Developed and deployed a real-time player tracking solution in Bird's Eye View, used live during the 2022 Asia Cup Cricket Tournament.

Stereo SLAM: Created 3D point clouds from stereo images and estimated motion/pose by applying the iterative Perspective-n-Point (PnP) algorithm for 2D-3D correspondences.

Pose Graph Optimization: Implemented the Levenberg–Marquardt algorithm to optimize robot poses using Odometry and Loop Closure constraints for 1D and 2D Simultaneous Localization and Mapping (SLAM).

Unity Game for Amblyopia: Developed a Unity-based game for diagnosing Amblyopia, controlled through eye gaze movements tracked by an eye-tracking device.

Neuro Rehab Systems: Designed a rehabilitation tool aimed at assisting recovery from nervous system injuries and minimizing functional impairments.

TECHNICAL STRENGTHS

Languages: Python, Java, C++, C#, Node JS, Javascript

Frameworks: PyTorch, Keras, Tensorflow, OpenCV

Tools: CARLA, Unity 3D, TensorRT, Open3D, Matlab

RELEVANT COURSES

ML/AI Courses: Statistical Methods in AI, Computer Vision, Mobile Robotics, Topics in Optimization Methods, Topics in ML, Cognitive Science and AI

Core Science: Computer Programming, Operating Systems and Algorithms, Data Structures

Other Courses: Digital Image Processing, Digital Signal Processing, Linear Algebra, Probability and Random Processes, Discrete Mathematics

ACADEMIC SERVICE

Reviewer: NeurIPS (2025, 2026); CVPR (2024); WACV (2024, 2025); ACL (2024); ECCV (2024, 2026); ICRA (2024).

ACHIEVEMENTS

Morocco Solidarity Hackathon: Member of the winning team; developed a solution to predict human trafficking behavior from social media posts.

Qualcomm Innovation Fellowship: Led the winning team of the Qualcomm Innovation Fellowship (QIF) 2020 in India.

JEE Mains: Ranked in the top 0.2% nationally among 1,200,000 candidates in JEE Mains.

JEE Advanced: Secured a rank of 4539 among 150,000 candidates in JEE Advanced.

REFERENCES

Dr. Vineet Gandhi, Dr. K. Madhava Krishna